

Databases

The Relational Data Model

Dr Konrad Dabrowski

konrad.dabrowski@durham.ac.uk

Online Office Hour:

Mondays 13:30–14:30

See Duo for the Zoom link

Relational Data Model

- Data Definition Language (DDL) is **too low-level**
 - not easily understandable by most of users

⇒ We need a Data Model:

- a collection of **intuitive** concepts describing *data*, their *relationships* and *constraints*
- Relational Data Model:
 - **relations** between data → stored in **tables**
 - based on the concept of **mathematical relations**
 - the most widely used Data Model (for **structured** data)
- Intuitively:
 - in our data, every **entity** (e.g. customer) combines (or “**relates**”) various “**attributes**” together (e.g. *name*, *id*, *address*, etc.)

Relational Data Model

- The **schema** of a relation:
 - the description of a particular collection of data in the model
- Let A_1, A_2, \dots, A_n be a set of **attributes** that can be “related”
 - i.e. there exists an entity with some values for these attributes
- Then $R(A_1, A_2, \dots, A_n)$ is the **schema** of the relation R
e.g. if R denotes **customers**, the schema of R could be:
 $R(\text{name}, \text{id}, \text{address}, \text{town}, \text{date-of-last-purchase})$
 - *this means: every customer has values for exactly these attributes*
- In a relation schema:
 - the **ordering** of the **attributes** does **not** matter!
(in contrast to mathematical relations)

Relational Data Model

- A database has many **entities**
 - each with its own **attributes**
 - This information is decomposed into smaller pieces
 - every **relation** stores only one piece of the information
 - Example:
 - the relation **account** stores info for the accounts
 - **depositor** stores info about which customer has this account
 - **customer** stores info about all customers
 - the whole info of the DB could be stored in **one** relation, e.g.
 $R(\text{account}, \text{balance}, \text{cust-name}, \dots)$
- but then:
- **data duplication** (e.g. two customers have the same account)
 - we need many **null** values (e.g. insert customer with no account)

Relational Data Model

Cust-id	Cust-Name	Cust-address	Cust-city
192-83-7465	Johnson	Bakerstr. 1	Durham
019-28-3746	Dunn	Halkinstr. 23	Newcastle
677-89-9011	Richardson	Pensburyst. 12	London
182-73-6091	Edwards	Gardenstr. 25	Manchester

Customer

Account

customer with
no accounts

customer with
3 accounts

Account number	Balance
A-101	500
A-215	700
A-102	400
A-305	350
A-222	700

Cust-id	Account number
192-83-7465	A-101
019-28-3746	A-215
677-89-9011	A-102
019-28-3746	A-305
019-28-3746	A-222

Depositor
(i.e. customer
with an account)

Relational Model Terminology

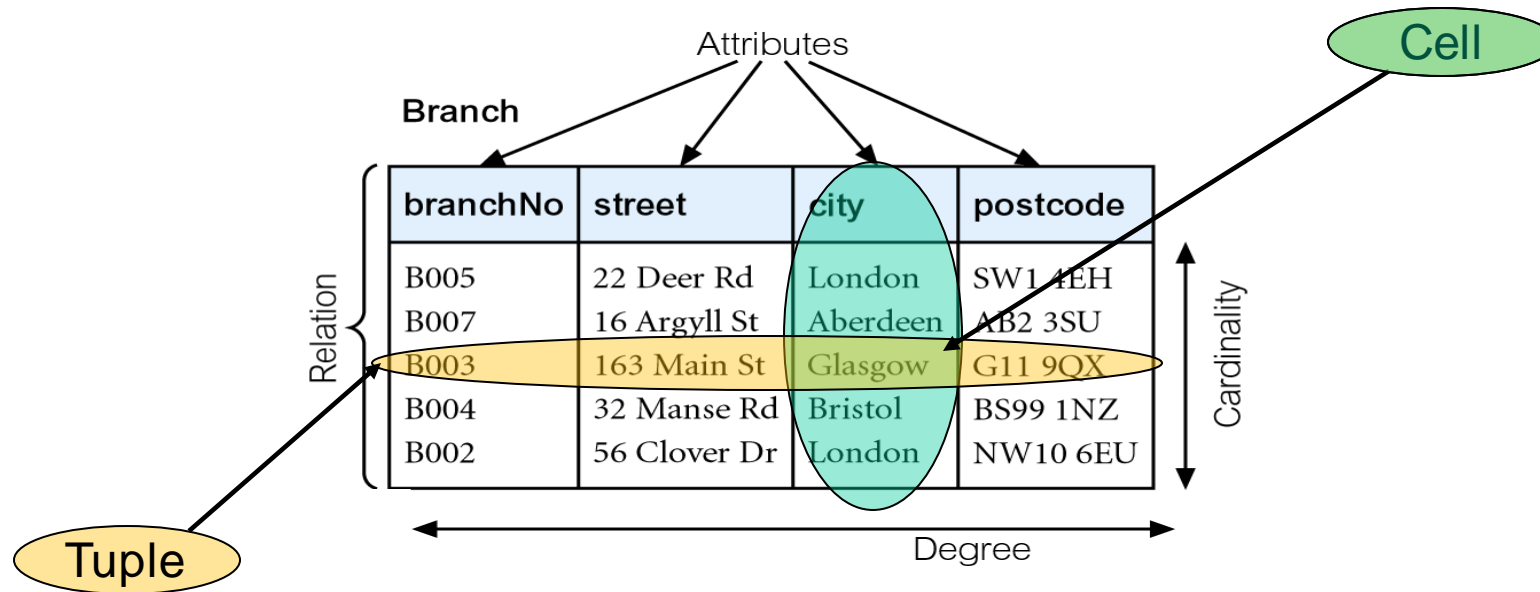
- A **relation** is a table (with **rows** and **columns**)
 - this only refers to the logical structure of the DB, not the physical structure
- An **attribute** is a named **column** of a relation
 - every attribute has a unique name
- The **domain** of an **attribute** is the set of allowable values
- A **tuple** is a **row** of a relation
 - every tuple has a concrete value for every attribute!
- A **cell** of a relation is the intersection of a row and a column
- The **degree** of a relation is the **number of attributes**
 - i.e. every row stores as many values as the degree of the relation
- The **cardinality** of a relation is the **number of tuples**

Relational Model Terminology

- A relation is **normalized** if it is “appropriately structured”, e.g.
 - every cell has exactly one value (not more / less!)
 - no repetitions of two identical rows
- A **Relational Database** is a **collection** of **normalized relations**
 - each of them with distinct relation names
- An alternative terminology for the Relational Model:

Formal terms	Alternative 1	Alternative 2
Relation	Table	File
Tuple	Row	Record
Attribute	Column	Field

Instances of Branch and Staff Relations



Staff

staffNo	fName	lName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

Relation

Instances of Branch and Staff Relations

- An illustration of the **domains** of the attributes:

Attribute	Domain Name	Meaning	Domain Definition
branchNo	BranchNumbers	The set of all possible branch numbers	character: size 4, range B001–B999
street	StreetNames	The set of all street names in Britain	character: size 25
city	CityNames	The set of all city names in Britain	character: size 15
postcode	Postcodes	The set of all postcodes in Britain	character: size 8
sex	Sex	The sex of a person	character: size 1, value M or F
DOB	DatesOfBirth	Possible values of staff birth dates	date, range from 1-Jan-20, format dd-mmm-yy
salary	Salaries	Possible values of staff salaries	monetary: 7 digits, range 6000.00–40000.00

- NULL** value:
 - a special case of a cell entry
 - it represents an **attribute value** that is:
 - either currently **unknown**
 - or **not applicable**
 - not the same as the value “0”!
 - may / may not belong to the domain of the attribute

Instances of Branch and Staff Relations

Another example for the domains:

- the **schema** of the relation R of the **customers** is:

$R(\text{name}, \text{address}, \text{town})$

- where:
 - the domain of “name” is {Johnson, Dunn, Edwards, Richardson}
 - the domain of “address” is {Bakerstr, Halkinstr, Gardenstr}
 - the domain of “town” is {London, Durham, Glasgow}

which one is a correct relation?



name	address	town
Johnson	Bakerstr	London
Edwards	Halkinstr	Durham
Richardson	Gardenstr	London



name	address	town
Johnson	Bakerstr	London
Edwards	Halkinstr	Durham
Richardson	Gardenstr	Leeds

Properties of Relations

- The **relation name** is **distinct** from all other relation names in the relational schema
- Each **attribute** (within a relation) has a **distinct** name
 - possibly two attributes of different relations may have the same name
 - e.g. “name”, “id” etc.
- Values of an **attribute** are all from the same **domain**
- Each **cell** of relation contains exactly **one** atomic (single) value
- Each **tuple** is **distinct** among the tuples of the relation
 - there are **no duplicate tuples**
- The **ordering** of **attributes** has **no significance**
 - unlike mathematical relations
 - e.g. Cartesian product of two sets: ordered pairs of elements
- The **ordering** of **tuples** has **no significance**
 - just unordered rows

Structuring concept: Keys

- How do we **uniquely** identify a **tuple** in a (normalized) table?
 - attribute **names** are **unique** within a table ('id', 'sex', ...)
 - **but** two **tuples** may **share** attribute **values** (both have sex 'M')
- Every table **must** have some **attributes**, such that:
 - their value **uniquely determines** a **tuple** of the table
 - these attributes are the **primary key** of the table

CD table

<u>CD #</u>	Song	Artist	Position	Month	Year
0001	Mandolin Wind	Rod Stewart	1	02	1974
0002	Gallows Pole	Led Zeppelin	2	04	1985
0003	Comfortably Numb	Pink Floyd	5	04	1989
0004	Paint it Black	Rolling Stones	6	09	1967

Structuring concept: Keys

- Candidate key: (of a relation)
 - a **minimal** (not minimum!) set of **attributes** (“keys”) whose values **uniquely identify** the **tuples**
- Primary key:
 - The candidate key selected to **identify** rows **uniquely** within the table
- Alternate key:
 - Those candidate key(s) not selected as primary key
- Simple key:
 - The key consists of only one attribute
- Composite key:
 - The key consists of several attributes

Primary keys

- An example:

Employee				
Initial	Surname	Job Title	Car No.	Department
N	Cook	Salesman	5	Sales
A	Randell	Programmer	null	Computer
M	Cook	Consultant	12	Insurance

- Is it possible to use *Surname* as the primary key?

Maybe sometimes, but not in this case ...

- What about using *Initial*? Even worse!

- In this example, initials are all different, so *Initial* can be a primary key
- However: if *Mark Washington* gets a job at this company, the DBMS will not allow his name to be added

Primary keys

- An example:

Employee				
Initial	Surname	Job Title	Car No.	Department
N	Cook	Salesman	5	Sales
A	Randell	Programmer	null	Computer
M	Cook	Consultant	12	Insurance

- Could we use *Initial* and *Surname* as a (composite) primary key?
 - Much better, but still...
 - *What happens if another Nick Cook appears?*
- Which primary key should be used?
 - The best choice is the *Employee No.* which is unique for every person!
 - Or alternatively: to add another new “*identification*” attribute

Composite keys

- We can create **composite keys**:
 - by combining two (or more) attributes
- In the table below *Song* and *Artist* are a composite key
 - this key uniquely identifies a row

<u>Song</u>	<u>Artist</u>	Position	Month	Year
Tubular Bells	Mike Oldfield	1	02	1974
Gallows Pole	Led Zeppelin	2	04	1985
Comfortably Numb	Pink Floyd	5	04	1989
Paint it Black	Rolling Stones	6	09	1967

- Why use composite keys?
 - Providing that an artist never reissues the same song title, *Song* and *Artist* are a valid primary key
 - *Song* on its own would run into difficulties (other artists may use this title)

Foreign keys

- To keep track of who borrowed a CD, we need to store:
 - the details of our **customers**
 - the details of all the **transactions**

Customer table

<u>Cust #</u>	Name	Address
001	Konrad Dabrowski	Stockton Road, Durham
002	Janet Lavery	Stockton Road, Durham
003	Stephan Jamieson	Old Elvet, Durham

Transaction table

<u>Trans#</u>	CD#	Cust #	Hire date	Duration
00001	0001	001	12/9/06	1
00002	0001	002	14/9/06	3
00003	0004	002	15/9/06	1

- How do we “match” the customers in the two tables?
 - using **foreign keys!**

Foreign keys

- Foreign key:
 - An **attribute** in one **table A** whose values must:
 - either **match** the **primary key** of another **table B** (then **A** references **B**)
 - or be **NULL** (e.g. staff has not been yet assigned to a branch)

table B:

Diagram illustrating the structure of table B:

- Attributes**: branchNo, street, city, postcode
- Branch**: points to the branchNo attribute.
- Relation**: points to the entire table structure.
- Cardinality**: points to the vertical axis of the table.
- Degree**: points to the horizontal axis of the table.

branchNo	street	city	postcode
B005	22 Deer Rd	London	SW1 4EH
B007	16 Argyll St	Aberdeen	AB2 3SU
B003	163 Main St	Glasgow	G11 9QX
B004	32 Manse Rd	Bristol	BS99 1NZ
B002	56 Clover Dr	London	NW10 6EU

table A:

Diagram illustrating the structure of table A:

- Staff**: points to the entire table structure.
- Relation**: points to the entire table structure.

staffNo	fName	lName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

Integrity Constraints

- So far we have seen:
 - domain constraints for the attributes
- Entity integrity:
 - every attribute of a primary key cannot be null
(otherwise we do not need all attributes of the primary key to identify the tuples)
- Purpose of entity integrity:
 - guarantees that each entity has unique identifier
 - ensures that foreign key values can reference primary key values
- Example:
 - no invoice can have a duplicate number, nor can it be null
 - all invoices uniquely identified by the invoice number

Integrity Constraints

- **Referential integrity:**
 - a **foreign key** either **matches** the **primary key** in the table it refers to
 - or it is **null**
- **Purpose of referential integrity:**
 - any reference between tables is valid (or it has not been set yet)
 - prevents deleting a row in a table B , if the primary key of B has a matching foreign key in another table A
- **Example:**
 - a customer will be always assigned to a valid sales representative
 - unless (s)he is not yet assigned to any representative

Summary: Characteristics of a relational table

- A relation is represented by a two-dimensional table
- Each row (tuple) signifies an entity occurrence
- No two rows can be identical (each row of the table is unique)
- Each column represents an attribute and has a distinct name
- The intersection of a row and a column has a single value (atomic)
- All values in a column must be of the same type (e.g. integers)
 - they have the same domain
- One (or more) attributes uniquely identify each row (primary key)
 - primary keys are not allowed a NULL value
- Two tables can be dependent
 - the primary key is the foreign key of another table
- The ordering of rows and columns does NOT matter

Views

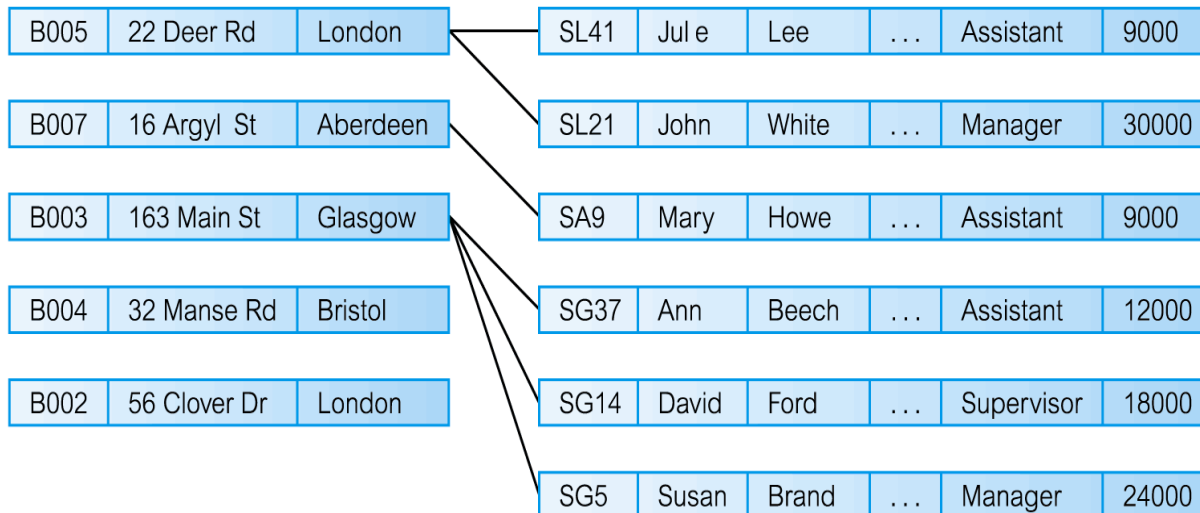
- So far all relations (tables) we have seen:
 - base relations
 - its tuples are physically stored in the database
- A different type of a relation: a view
 - a virtual relation
 - it does not exist physically in the database
- The content of a view:
 - is derived from one (or more) base relations
 - is computed upon request by a user, at the time of request
 - changes when the underlying base relations change
- Main use:
 - show customized information to every user
(e.g. show “loan number” but not “amount borrowed”)
 - compute dynamic quantities (e.g. “age” from “date-of-birth”)

Alternatives to the Relational Data Model

Other models before the development of the Relational Data Model:

1. Network Data Model

- records (tuples) appear as **nodes**
- relationships (foreign keys) appear as **edges**



Branch

<u>branchNo</u>	street	city	postCode
B005	22 Deer Rd	London	SW1 4EH
B007	16 Argyll St	Aberdeen	AB2 3SU
B003	163 Main St	Glasgow	G11 9QX
B004	32 Manse Rd	Bristol	BS99 1NZ
B002	56 Clover Dr	London	NW10 6EU

Staff

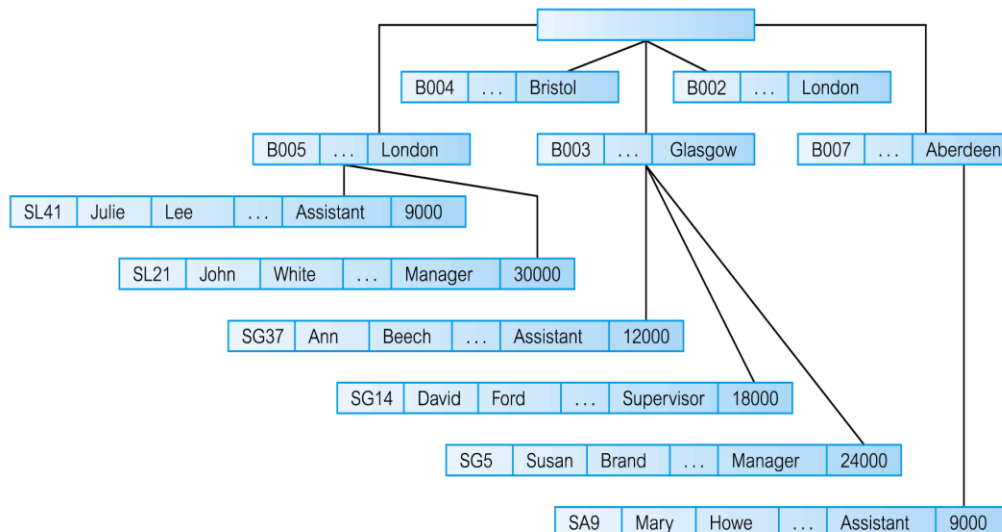
<u>staffNo</u>	fName	lName	position	sex	DOB	salary	<u>branchNo</u>
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

Alternatives to the Relational Data Model

Other models before the development of the Relational Data Model:

2. Hierarchical Data Model

- special case of the Network Data Model, where the graph is a **tree graph**
- its structure mirrors **parent-child relationship** (**one parent, many children**)
- limitations of this model, e.g.
 - deleting a parent,
 - adding a record without a parent



Branch

branchNo	street	city	postCode
B005	22 Deer Rd	London	SW1 4EH
B007	16 Argyll St	Aberdeen	AB2 3SU
B003	163 Main St	Glasgow	G11 9QX
B004	32 Manse Rd	Bristol	BS99 1NZ
B002	56 Clover Dr	London	NW10 6EU

Staff

staffNo	fName	lName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

Additional Slides

Mathematical Relations

In mathematics (background):

- for n sets D_1, D_2, \dots, D_n , their *Cartesian product* is:

$$D_1 \times D_2 \times \dots \times D_n = \{(d_1, d_2, \dots, d_n) \mid d_1 \in D_1, \dots, d_n \in D_n\},$$

i.e. all the *ordered* n -tuples of these sets

- example: $D_1 = \{1, 2\}$, $D_2 = \{a, b\}$, $D_3 = \{x, y\}$, thus:

$$\begin{aligned} D_1 \times D_2 \times D_3 = \{ & (1, a, x), (1, a, y), \\ & (1, b, x), (1, b, y), \\ & (2, a, x), (2, a, y), \\ & (2, b, x), (2, b, y) \} \end{aligned}$$

- in mathematics, a *relation* among the sets D_1, D_2, \dots, D_n is *any (ordered) subset* $R \subseteq D_1 \times D_2 \times \dots \times D_n$

Mathematical Relations

However,

in the **Databases** terminology,

the **order of attributes** (D_1, D_2, \dots, D_n) ,

i.e. **column headings**, in a relation does **not** matter!

- in mathematics, a **relation** among the sets D_1, D_2, \dots, D_n is **any** (**ordered**) **subset** $R \subseteq D_1 \times D_2 \times \dots \times D_n$