

Databases

Distributed Architectures and DBMS

Dr Konrad Dabrowski

konrad.dabrowski@durham.ac.uk

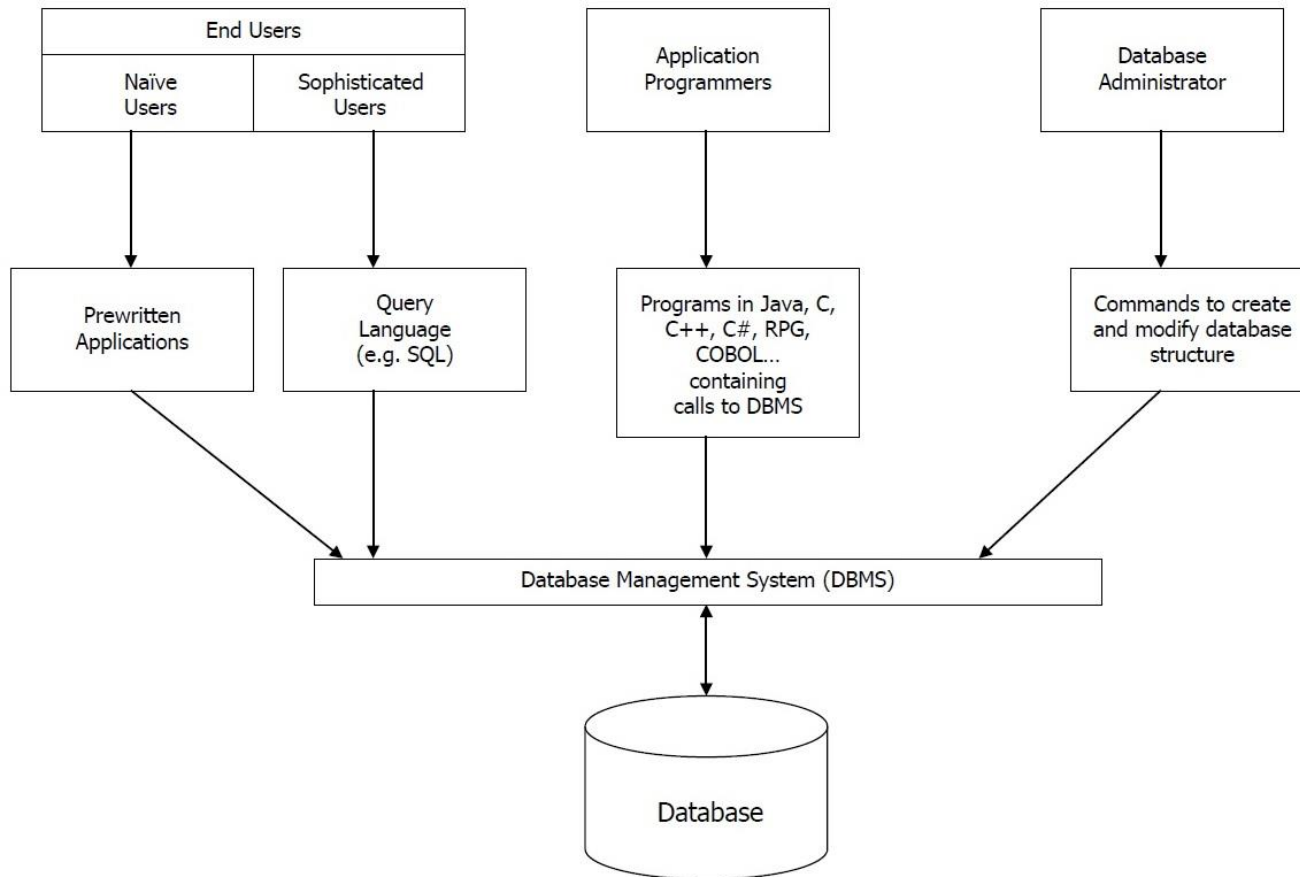
Online Office Hour:

Mondays 13:30–14:30

See Duo for the Zoom link

Multi-user DBMS Architectures

- So far we have seen this model of interaction between end-users and the database:
 - one central DBMS
 - users interact with DBMS using an application program




Multi-user DBMS Architectures

- Several issues still need clarification:
 - is the DBMS on the end-user's computer?
 - how are the users connected to the DBMS?
(e.g. physically, through LAN etc.)
 - do we have one or more DBMS?
 - which computations are performed where?
(e.g. on the user's computer, on some server etc.)
 - is the database stored in one or many places?
- The answers are not unique:
 - different company needs require different approaches

Multi-user DBMS Architectures

- **Teleprocessing Architecture:**

- the traditional (and most basic) architecture
- **one computer** with a single Central Processing Unit (CPU)
- **many** (end-user) **terminals**,
all cabled to the central computer
- a terminal sends messages to the central computer (through the user's application program)
- all data processing in the central computer:
 - receive messages from terminal
 - process the queries
 - send the results back to the terminal



Multi-user DBMS Architectures

- **Teleprocessing Architecture:**


- tremendous burden on the central computer
⇒ decreased performance

- Nowadays, the trend is towards **downsizing:**

- replace expensive mainframe computers with cost-effective networks of personal computers
- achieve the same / better results


- The next two architectures:

- **file-server** architecture
- **client-server** architecture



File-Server architecture

- Processing is distributed around a computer network
 - typically through a Local Area Network (LAN)
 - one **central file-server**
(computer containing the database)
 - every **workstation** has its **own DBMS** and its **own user application**
 - workstations **request files** they need from the file server
 - file server acts like a “shared hard disk” (it has no DBMS!)
- For example, the user request is:
 - “find the names of staff who work in the branch at 163 Main St.”
 - the user application creates an SQL query for this

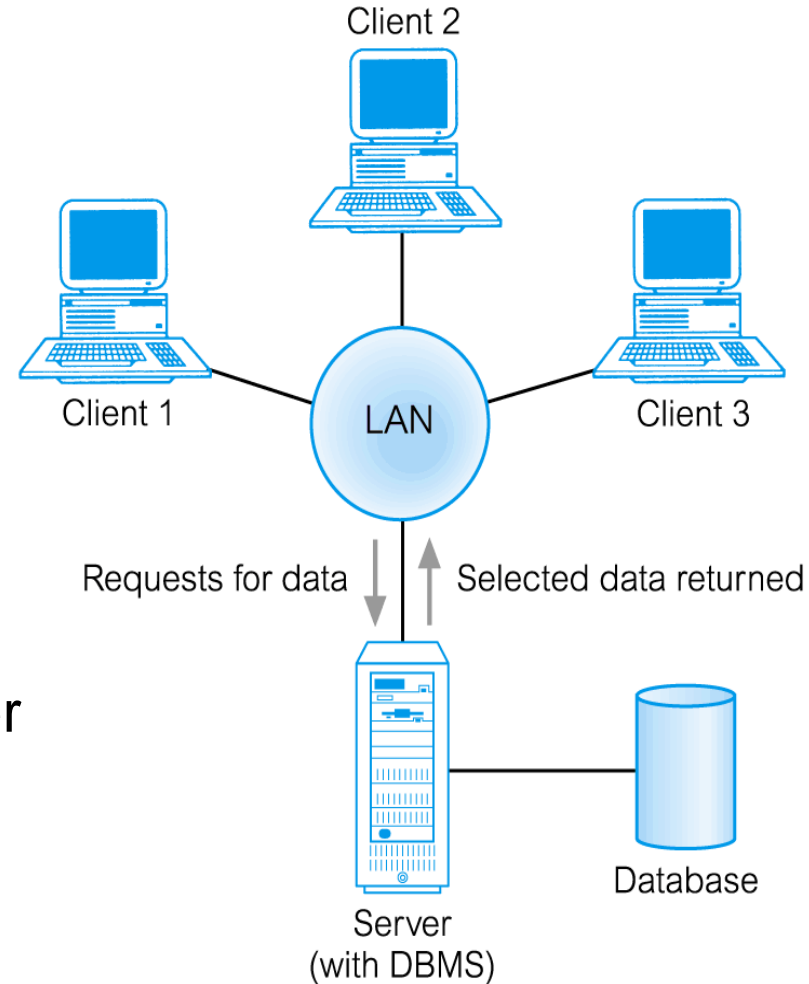


File-Server architecture

- The SQL query:
SELECT fName, lName
FROM Branch b, Staff s
WHERE b.branchNo = s.branchNo **AND** b.street = '163 Main St.'
- The file-server has no knowledge of SQL (has no DBMS!)
 - the **user's DBMS** has to request the **whole tables** Branch, Staff
- Therefore:
 - very large amount of **network traffic** (the tables may be huge)
 - a **full copy** of the **DBMS** is required on each workstation
 - **concurrency / recovery / integrity control** is more difficult, since multiple DBMSs access the same files simultaneously
- The solution to these problems:
 - **client-server** architectures

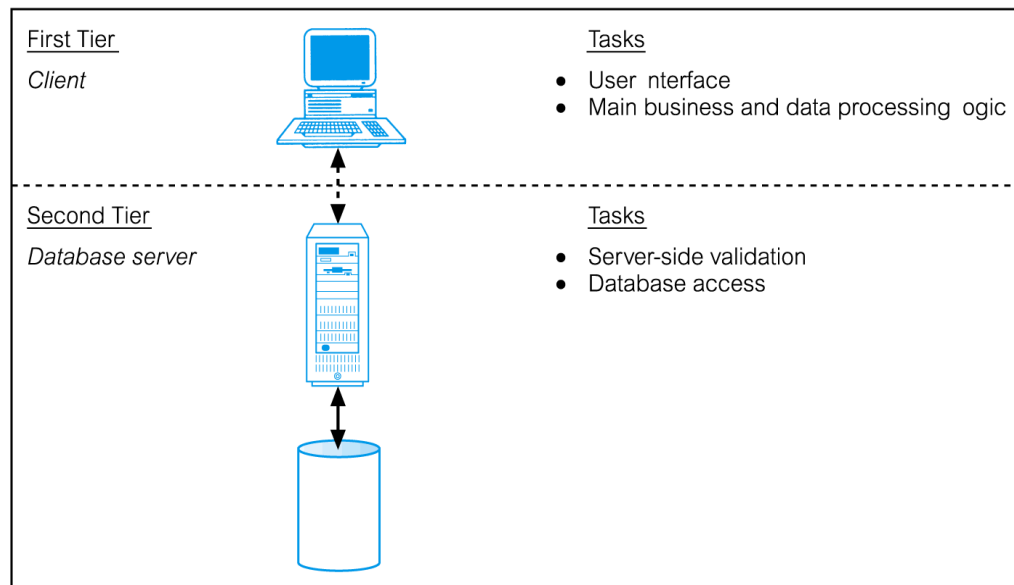
Client-Server architecture

- **Client:** requires some resource
- **Server:** provides the resource
- Client / Server are not always in the same machine / place
- **Two-tier architecture:**
 - **Tier 1 (client):** responsible for the **presentation** of data to the user
 - **Tier 2 (server):** responsible for supplying **data services** to the user
- Typical procedure:
 - **user** gives a request to the **client**
 - **client** generates **SQL query** and sends it to the **server**
 - **server** accepts, processes the query and sends the result to the **client**
 - **client** formats the result for the **user**



Client-Server architecture

- Many advantages:
 - **increased performance**: many client CPUs work in parallel
 - **reduced hardware costs**: only the server needs increased storage and computational power
 - **reduced communication costs**: less data traffic (no unnecessary data transmitted)
- Database is still centralized:
 - not a distributed database!

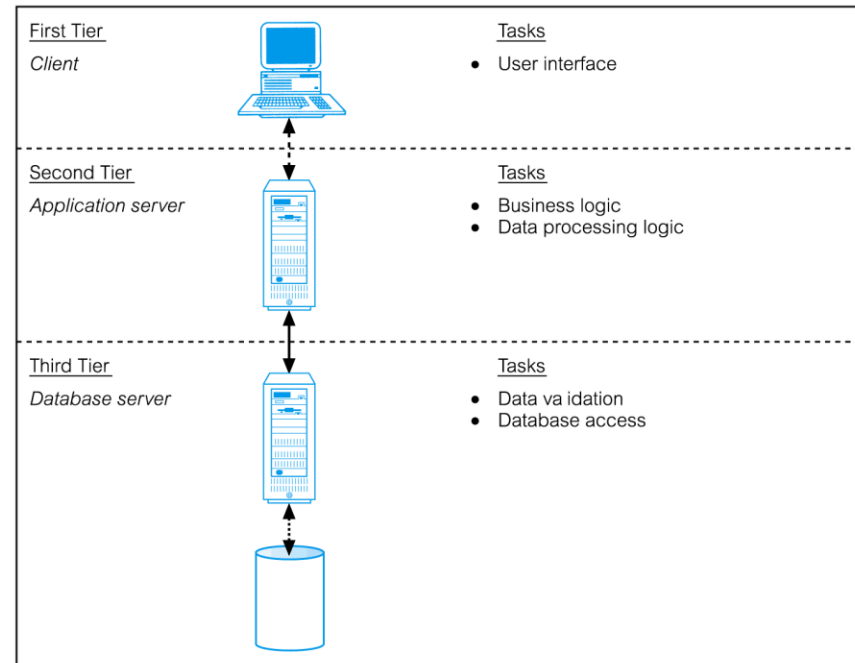


Three-Tier Client-Server Architecture

- In modern systems: 100s / 1000s of users
⇒ need for increased enterprise scalability
- Main problem of the client that prevent scalability:
 - a “fat client” (many users) requires extensive resources on disk space / RAM / CPU power
- A new variation of client-server architecture (1995):
 - three layers, potentially running on different platforms
 - First tier: user-interface layer (on end-user’s computer)
 - Second tier: application server (connects to many users)
 - Third tier: database server
(contains DBMS, communicates with the application server)
 - “thin clients” ⇒ increased performance of user’s computer

Three-Tier Client-Server Architecture

- Best example for a client: **internet browser** (fast & light)
- Advantages:
 - **smaller hardware cost** for the “thin clients”
 - easier application **maintenance** (centralized in one tier)
 - **easier to modify/replace** one tier without affecting the others
 - **easier load balancing** between the different tiers
 - maps naturally to **Web applications**
- It can be extended:
 - separation of tasks into *n* intermediate tiers
 - ⇒ increased flexibility / scalability



Distributed DBMS


- So far we have seen **centralized database systems**:
 - single database, located at one site
 - controlled by one DBMS
 - We can improve database performance:
 - using **networks of computers** (**decentralized** approach)
 - it mirrors the organizational structure:
 - logically distributed into divisions, departments, projects...
 - physically distributed into offices, units, factories...
 - Main targets:
 - make **all data** accessible to **all units** (avoid “islands of information”)
 - **store** the data **proximate** to the location
where it is **most frequently used**
- ⇒ full functionality and efficiency

Distributed DBMS

- **Distributed database:**
 - a logically interrelated collection of **shared data**, **physically distributed** over a **network**
- **Distributed DBMS (DDBMS):**
 - the **software system** that can manage the distributed database
 - it makes the distribution **transparent (invisible)** to users
- **In a DDBMS:**
 - a **single logical database**, which is **split** into **fragments**
 - each **fragment** is stored on one (or more) **computers**, under the control of a **separate DBMS**
 - all these computers are connected by a **communications network**
 - sites have **local autonomy**: **independent processing** of **local data** (via **local applications**)
 - sites have access to **global applications** (to process data fragments stored on other computers)

Distributed DBMS

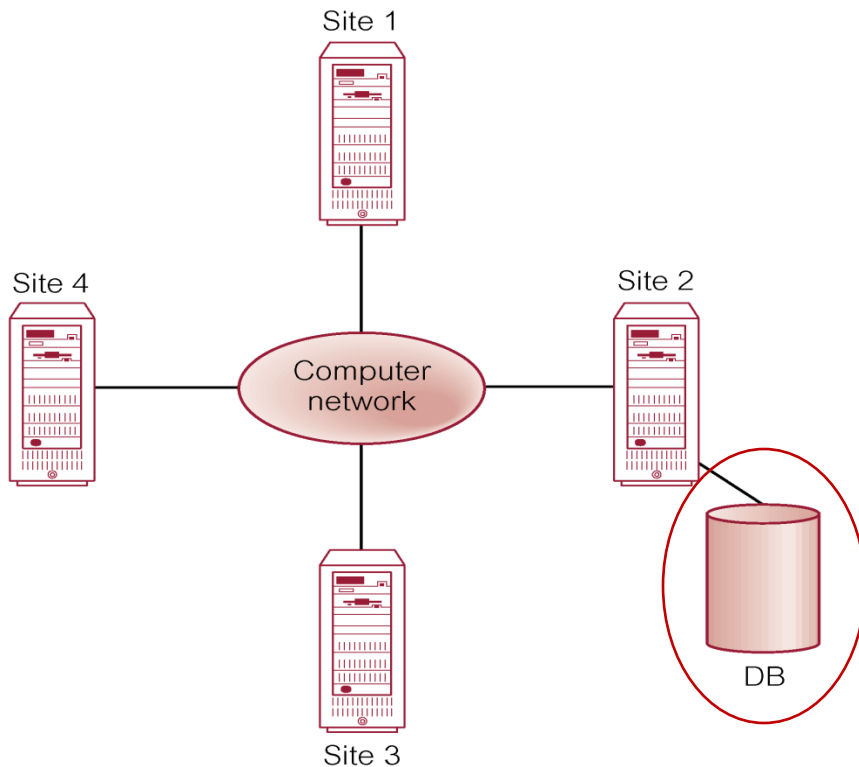
- Not all sites have local applications / local data
- All sites have access to global applications
- Data fragments may be replicated in more sites
 - data consistency must be considered
(same values for all replications)




Distributed Processing vs. Distributed DBMS

- **Distributed processing:**
 - a **centralized database** that is accessed over a **computer network**
 - for example: the client-server architecture
 - **not** the same as **distributed DBMS!**

distributed processing:



distributed DBMS:



Design of a distributed DBMS

- In addition to ER modelling, we also have to consider:
 - **Fragmentation**
 - how to break a relation into fragments
 - fragments can be horizontal / vertical / mixed
 - **Allocation**
 - how fragments are allocated at the several sites
 - aim is to reach an “optimal” distribution (efficient, reliable...)
 - **Replication**
 - which fragments are stored at multiple sites (& at which sites)
- Choices for Fragmentation and Allocation:
 - based on **how** the database is to be **used**
 - **quantitative** and **qualitative** information needed

Design of a distributed DBMS

- **Quantitative information:** (mainly for **fragmentation**)
 - the frequency with which specific transactions are run
 - the (usual) sites from which transactions are run
 - desired performance criteria for the transactions
- **Qualitative information:** (mainly for **allocation**)
 - the relations / attributes / tuples being accessed
 - the type of access (read / write)
- Strategic objectives for the choices about the fragments:
 - **locality of reference**
 - data to be stored close to where it is used
 - if a fragment is used at several sites → replication is useful
 - **reliability and availability**
 - improved by replication
 - if one site fails, there are other fragment copies available

Design of a distributed DBMS

- Further strategic objectives:
 - acceptable performance
 - bad allocation results in “bottleneck” effects
(a site receives too many requests \Rightarrow bad performance)
 - also: bad allocation causes underutilized resources
 - cost of storage capacities
 - cheap mass storage to be used at sites, whenever possible
 - this must be balanced against locality of reference
 - minimal communication costs
 - min. retrieval costs when max. locality of reference,
or when each site has its own copy of data
 - but when replicated data are updated:
 - \rightarrow all copies of this data must be updated
 - \rightarrow increased network traffic / communication costs

Design of a distributed DBMS

- Four alternative strategies for the placement of data:
 - centralized
 - single database and DBMS
 - stored at one site with users distributed across the network
 - not distributed!
 - partitioned
 - database partitioned into disjoint fragments
 - each data item assigned to exactly one site (no replication!)
 - complete replication
 - complete copy of the database at each site
 - selective replication
 - combination of partitioning, replication, and centralization

Design of a distributed DBMS

- Four alternative strategies for the placement of data:

Balance of the strategic objectives

	Locality of reference	Reliability and availability	Performance	Storage costs	Communication costs
<u>Centralized</u>	Lowest	Lowest	Unsatisfactory	Lowest	Highest
<u>Fragmented</u>	High ^a	Low for item; high for system	Satisfactory ^a	Lowest	Low ^a
<u>Complete replication</u>	Highest	Highest	Best for read	Highest	High for update; low for read
<u>Selective replication</u>	High ^a	Low for item; high for system	Satisfactory ^a	Average	Low ^a

^a Indicates subject to good design.

Design of a distributed DBMS

Three **correctness rules** for the **partitioned placement** (i.e. with no replication):

1. Completeness

- if relation R is decomposed into fragments R_1, R_2, \dots, R_n , each data item that can be found in R must appear in at least one fragment R_i

2. Reconstruction

- it must be possible to define a rel. algebra expression that can reconstruct R from its fragments

3. Disjointness

- if a data item appears in fragment R_i , it should not appear in any other fragment
- **exception** for vertical fragmentation: **primary key** attributes must be **repeated** for the **reconstruction** !

Design of a distributed DBMS

Three main types of fragmentation:


1. horizontal

- a subset of the tuples of the relation





2. vertical

- a subset of the attributes of the relation



3. mixed

- a vertical fragment that is then horizontally fragmented
- or a horizontal fragment that is then vertically fragmented



Horizontal fragmentation

- Assume there exist two property types: 'Flat' and 'House'
- We have a relation R with all properties for rent
- The **horizontal fragmentation** of R (by property-type) is:
 - $P_1 = \sigma_{\text{type}='House'}(\text{PropertyForRent})$
 - $P_2 = \sigma_{\text{type}='Flat'}(\text{PropertyForRent})$
- This fragmentation may be **useful**:
 - e.g. if we have separate applications dealing with flats / houses
- and it is **correct**:
 - **completeness**: each tuple is either in P_1 or in P_2
 - **reconstruction**: R can be reconstructed from the fragments P_1, P_2
$$R = P_1 \cup P_2$$
 - **disjointness**: there is no property that is both 'flat' and 'house'

Relational Algebra operator " σ ":
"select" specific rows of a table

Vertical fragmentation

- For every staff member in a company:
 - the payroll department requires: staffNo, position, sex, salary
 - the HR department requires: staffNo, fName, DOB, branchNo
- We have a relation *Staff* with all staff members
- For this example, the vertical fragmentation of *Staff* is:

- $S_1 = \Pi_{\text{staffNo, position, sex, salary}}(\text{Staff})$

Relational Algebra operator “ Π ” :
“project” specific columns of a table

- $S_2 = \Pi_{\text{staffNo, Name, DOB, branchNo}}(\text{Staff})$

- Both fragments include the primary key staffNo
 - to allow reconstruction of *Staff* from S_1 and S_2
- This fragmentation is useful:
 - fragments are stored at the departments that are needed
 - performance for every department is improved
(as the fragment is smaller than the original relation *Staff*)

Vertical fragmentation

$$S_1 = \Pi_{\text{staffNo, position, sex, salary}}(\text{Staff})$$

$$S_2 = \Pi_{\text{staffNo, Name, DOB, branchNo}}(\text{Staff})$$

- This fragmentation is **correct**:
 - **completeness**:
 - the primary key staffNo belongs to both S_1 and S_2
 - each other attribute is either in S_1 or in S_2
 - **reconstruction**: R can be reconstructed from the fragments S_1 , S_2 using the Natural join operation:
$$\text{Staff} = S_1 \bowtie S_2$$
 - **disjointness**: the fragments are disjoint except the primary key (which is necessary for the reconstruction)

(Dis-)advantages of distributed DBMS

- **Advantages:**

- Reflects organizational structure
- Improves shareability and local autonomy
- Improved availability and reliability
- Improved performance
- Smaller hardware cost
- Scalability

- **Disadvantages:**

- Complexity
- Higher maintenance cost
- Security
- Integrity control more difficult
- Design more complex
- Lack of experience in the industry

