

Lecture 2: Process Management – Part 1

Barnaby Martin


`barnaby.d.martin@dur.ac.uk`

Process State


As a process executes, it changes state (thus status):

- **New** – the process is being created
- **Running** – instructions are being executed
- **Waiting** – the process is waiting for some event to occur
- **Ready** – the process is ready to be dispatched to the CPU.
- **Terminated** – – the process has completed its execution, or some other event causing termination.

Process State



CPU Bursts



Multiprogramming

Independent processes cannot affect or be affected by the execution of other processes.

Co-operating processes can affect or be affected by the execution of another process — there are advantages of implementing process co-operation.

Multiprogramming

The advantages of multiprogramming / implementing process co-operation may include:

- Computation speed-up: if there are multiple CPUs and/or cores.
- Convenience: single user wants to run many tasks.
- Information sharing: for example shared files.
- Modularity: programming (e.g. OOP)

Multiprogramming supports the co-operation of processes.

Multiprogramming

The aim of multiprogramming is to **maximize** CPU utilization.

With multiprogramming several processes are kept in memory concurrently.

When one process is **waiting**, the operating system executes (to **running**) one of the processes sitting in the **ready** queue.

CPU Scheduling is the method used to support multiprogramming in process management.

CPU Scheduling

CPU scheduling can be viewed as processes 'taking turns'.

It provides a basis of a fair and efficient sharing of system, in this case CPU, resources.

As the CPU is a primary computer resource therefore scheduling is fundamental to operating system design.

Schedulers

An operating system operates three types of scheduler: long-term, medium-term, and short-term (CPU):

The **Long-term scheduler (job scheduler)** selects which processes should be brought into the ready queue.

The **Medium-term scheduler** removes processes from active contention for the CPU by swapping processes in and out of the ready queue.

This temporarily reduces the number of processes that the short-term scheduler (CPU scheduler) must choose between.

CPU Scheduler

The **Short-term scheduler (CPU scheduler)** selects the process to be executed next and allocated to the CPU.

This decision is initiated when processes:

1. Switch from running to waiting state, or
2. Switch from running to ready state, or
3. Switch from waiting to ready, or
4. Terminate

CPU Scheduling - considerations

Design considerations for scheduling.

When the CPU switches to another process:

- The system must save the context of the old process.
- Load any saved context for the new process.

Context-switch time is an **overhead**: the system undertakes no user processing while switching.

Cost vs. Benefit.

Scheduling Criteria

- **CPU utilization**: keep the CPU busy!
- **Throughput**: number of processes that complete their execution within a specific number of time units.
- **Turnaround time**: amount of time to execute a particular process.
- **Waiting time**: total (accumulated) amount of time a process has been waiting in the ready queue.
- **Response time**: amount of time it takes from when a request was first submitted to the ready queue until the first response is produced (access to the CPU). NB: Not the time to complete the process!

Scheduling Algorithms

There are a number of different algorithms including:

- First Come, First Served (FCFS)
- Shortest Job First (SJF)
- Shortest Remaining Time First (SRTF)
- Priority (with or without pre-empting)
- Priority (with or without ageing)
- Round-Robin (RR)
- Multilevel queue / feedback queue

Algorithm Evaluation

Deterministic modelling:

Take a predetermined workload (case) and analyse the performance of each algorithm.

Simulation:

Complex model of the system and the way it is used. Beware of Bonini's paradox!

Post-implementation:

Examine the running operating system (a.k.a. 'live' testing).

First Come, First Served (FCFS)

- The first process to request the CPU is allocated the CPU until it is completed.

First Come, First Served (FCFS)

Example(s)

First Come, First Served (FCFS)

- The average waiting time may or may not be lengthy!
- A simple algorithm to implement.
- May result in a 'convoy' effect, for example short processes waiting on a long process to finish.

Shortest-Job-First (SJF)

- Compares each process waiting to determine the length of its next CPU burst.
- Use these lengths to schedule the process with the shortest time.
 - If two processes have the same length then FCFS.
- Non pre-emptive.
 - Once the CPU is given to the process it cannot be pre-empted until the process has completed its CPU burst.

Shortest-Job-First (SJF)

Example(s)

Shortest-Job-First (SJF)

- The average waiting time?
- A simple algorithm to implement?
- Long processes waiting on a short process to finish.
Worse?

Round Robin (RR)

Each process gets a small unit of CPU time: a time quantum or time slice usually $q = 10$ to 100 milliseconds.

After this time has elapsed, the process is pre-empted and added to the end of the ready queue.

If there are n processes in the ready queue and the time quantum is q , then each process gets $\frac{1}{n}$ of the CPU time in chunks of at most q time units at once.

Round Robin (RR)

Example(s)

Round Robin (RR)

- The average waiting time?
- A simple algorithm to implement?
- Short vs. long processes?